

Keep AI calls out of Convex mutations

Sanjay Suresh

March 2026

Convex mutations are transactional and meant to be deterministic — they can re-run, and the runtime assumes the same inputs produce the same writes. Anything non-deterministic or network-bound (an LLM call, token streaming, an embeddings request) breaks that assumption.

The right shape is: an **action** does the external call, then hands its result to a **mutation** that persists it. The action is where side effects live; the mutation stays a clean, replayable write.

I learned this the slow way on AIManager — chat responses were occasionally double-writing because the model call sat inside the mutation and got retried. Moving the call into an action made the whole thing predictable.