

Quantum Error Correction and Fault Tolerance

A synopsis of Gottesman's *An Introduction to Quantum Error Correction and Fault-Tolerant Quantum Computation*

Sanjay Suresh

June 2026

Abstract

Quantum states are delicate, so a useful quantum computer must protect fragile qubits with error correction and then compute on the encoded data without that protection unravelling. This synopsis follows D. Gottesman's review ([arXiv:0904.2557](https://arxiv.org/abs/0904.2557)): why quantum error correction looks impossible at first, the nine-qubit code that makes it possible, the conditions a code must satisfy, the stabilizer formalism and its links to classical coding theory, and then the harder problem of fault tolerance — computing reliably with noisy components — culminating in the threshold theorem.

1 Why it looks impossible

Three obstacles separate the quantum problem from the classical one.

No cloning. The classical repetition code copies a bit, $0 \rightarrow 000$, and corrects by majority vote. The quantum analogue $|\psi\rangle \rightarrow |\psi\rangle |\psi\rangle |\psi\rangle$ is forbidden: if such an operation existed, linearity would force $|\psi\rangle + |\phi\rangle \mapsto |\psi\rangle |\psi\rangle + |\phi\rangle |\phi\rangle$, contradicting the definition's $(|\psi\rangle + |\phi\rangle)(|\psi\rangle + |\phi\rangle)$ by the cross terms.

Continuous errors. An over-rotation sends $\alpha|0\rangle + \beta e^{i\phi}|1\rangle$ to $\alpha|0\rangle + \beta e^{i(\phi+\delta)}|1\rangle$, wrong by an arbitrarily small δ that accumulates. We will need to *digitize* such errors.

Decoherence. Measuring $\alpha|0\rangle + \beta|1\rangle$ collapses it; the environment constantly tries to “look.” Error correction must extract information about the error without learning anything about the encoded data.

2 The nine-qubit code and error digitization

Shor's code nests two repetition codes:

$$|\bar{0}\rangle = \frac{1}{\sqrt{2}}(|000\rangle + |111\rangle)^{\otimes 3}, \quad |\bar{1}\rangle = \frac{1}{\sqrt{2}}(|000\rangle - |111\rangle)^{\otimes 3}.$$

An arbitrary codeword $\alpha|\bar{0}\rangle + \beta|\bar{1}\rangle$ is a superposition, so no-cloning is not violated; the codewords span a linear coding subspace. The inner triplets correct bit flips by majority, the outer layer corrects phase flips by majority on the three signs, and parities are measured so as to reveal the error and nothing else.

The operators $I, X, Z, Y = iXZ$ generate the **Pauli group** $\mathcal{P}_n$: tensor products of I, X, Y, Z with a phase $\pm 1, \pm i$, so 4^{n+1} elements. Any two commute or anticommute, and $XZ = -ZX$. The *weight* $\text{wt}(Q)$ is the number of non-identity factors; the n -qubit Paulis span all $2^n \times 2^n$ matrices. That spanning property digitizes errors: a continuous phase error $R_{\theta/2} = \cos \frac{\theta}{2} I - i \sin \frac{\theta}{2} Z$ is corrected for free, because recording the error in an ancilla and measuring it projects onto either “no error” or “ Z error,” each of which the code handles.

Theorem (Linearity of correction). *If a code corrects errors A and B , it corrects any linear combination $\alpha A + \beta B$. In particular, a code correcting all weight- t Pauli errors corrects all errors on t qubits, including general decoherence $\rho \rightarrow \sum_i E_i \rho E_i^\dagger$.*

3 When does a subspace correct errors?

We must never confuse $|\bar{0}\rangle$ with $|\bar{1}\rangle$, even after errors. Distinguishing the error is *sufficient* but not *necessary* — the nine-qubit code cannot tell Z_1 from Z_2 , yet corrects either. This gives the **Knill–Laflamme conditions**.

Theorem (QECC condition). *A subspace C corrects a linear space of errors $\mathcal{E}$ iff for all $E \in \mathcal{E}$, $\langle \psi | E^\dagger E | \psi \rangle = C(E)$ independent of $|\psi\rangle$; equivalently $\langle \psi_i | E_a^\dagger E_b | \psi_j \rangle = C_{ab} \delta_{ij}$.*

If C_{ab} has maximal rank the code is **nondegenerate**, otherwise (as for Shor’s code) **degenerate**. A code correcting t errors has **distance** $d = 2t + 1$, and an $[[n, k, d]]$ code uses n physical qubits for k logical at distance d . Counting states bounds the parameters:

$$\left(\sum_{j=0}^t 3^j \binom{n}{j} \right) 2^k \leq 2^n \quad (\text{quantum Hamming, nondegenerate}), \quad n - k \geq 2d - 2 \quad (\text{Singleton}).$$

The five-qubit $[[5, 1, 3]]$ code saturates the Singleton bound (smallest correcting one error); the nine-qubit code is $[[9, 1, 3]]$.

4 The stabilizer formalism

The corrective parity checks are measurements of commuting Paulis such as $Z_1 Z_2$. The group they generate is the **stabilizer** S , an Abelian subgroup of $\mathcal{P}_n$ with $-1 \notin S$ and $M|\psi\rangle = |\psi\rangle$ for all $M \in S$, $|\psi\rangle \in C$. If an error E anticommutes with a generator M , then

$$M(E|\psi\rangle) = -EM|\psi\rangle = -E|\psi\rangle,$$

flipping the eigenvalue to -1 . Measuring the generators gives the **error syndrome**, which identifies E up to degeneracy $E^\dagger F \in S$.

Theorem (Stabilizer code). *If S is an Abelian subgroup of order 2^a of $\mathcal{P}_n$ with $-1 \notin S$, and d is the minimum weight in $N(S) \setminus S$ (where $N(S)$ is the normalizer), then the states fixed by S form an $[[n, n - a, d]]$ code.*

In the **binary symplectic representation** each Pauli is a pair $(a | b)$; two operators commute iff $a_1 \cdot b_2 + b_1 \cdot a_2 = 0$, and the stabilizer condition becomes self-orthogonality under this symplectic form.

5 A zoo of codes

CSS codes. Two classical codes with $C_2^\perp \subseteq C_1$ give an $[[n, k_1 + k_2 - n, \min(d_1, d_2)]]$ code (Calderbank–Shor–Steane), with codewords

$$|\bar{u}\rangle = \sum_{w \in C_2^\perp} |u + w\rangle, \quad u \in C_1,$$

so Z -parities are checked in the standard basis and X -parities in the Hadamard basis. The seven-qubit **Steane code** $[[7, 1, 3]]$ uses the Hamming code for both.

Codes over $\text{GF}(4)$. Identifying I, Z, X, Y with $0, 1, \omega, \omega^2$ turns commutativity into trace-Hermitian self-orthogonality of a classical $\text{GF}(4)$ code; the five-qubit code is a Hamming code over $\text{GF}(4)$.

Beyond stabilizers. Qudit codes, subsystem (operator) codes, decoherence-free subspaces (fully degenerate, passive, but only for special noise), and dynamical decoupling (fast pulses, no extra qubits) extend the toolbox; in practice one layers them.

6 Logical operators and the Clifford group

The normalizer modulo the stabilizer is the **logical Pauli group** $N(S)/S \cong \mathcal{P}_k$: $E \in N(S)$ preserves the code space, and if $E \notin S$ it acts as a logical $\bar{X}, \bar{Z}$. The **Clifford group** $\mathcal{C}_n = \{U : U\mathcal{P}_n U^\dagger = \mathcal{P}_n\}$ is generated by $H, P = \text{diag}(1, i)$, and CNOT, with e.g. $H : X \mapsto Z, Z \mapsto X$, and $\mathcal{C}_n/\mathcal{P}_n' \cong \text{Sp}(2n, \mathbb{Z}_2)$. By the **Gottesman–Knill theorem** Clifford circuits on stabilizer states are efficiently classically simulable, so Cliffords alone are not universal.

7 Fault tolerance: the idea

Encoding is not enough; we must compute and correct with noisy gates. A protocol is **fault tolerant** if a few component failures produce only a correctable number of errors. The danger is **error propagation**: even a perfect gate spreads errors,

$$UE|\psi\rangle = (UEU^\dagger)U|\psi\rangle, \quad \text{CNOT} : X \otimes I \mapsto X \otimes X, \quad I \otimes Z \mapsto Z \otimes Z,$$

so one error can become two. Rigorously (Aliferis–Gottesman–Preskill), correctness uses an *r-filter* (projector onto $Q|\psi\rangle$ with $\text{wt}(Q) \leq r$) and an *ideal decoder*: a gadget is fault tolerant if, whenever input errors plus gadget faults total at most t , (A) errors do not propagate too far within a block and (B) it performs the correct logical operation. The universal set is $\{H, \text{CNOT}, R_{\pi/8}\}$ with $R_{\pi/8} = \text{diag}(1, e^{i\pi/4})$.

8 Transversal gates and their limit

A **transversal gate** $U = \bigotimes_i U_i$ couples the i -th qubit of one block only to the i -th of another, so it propagates an error to at most one qubit per block — automatically fault tolerant. Logical Paulis are transversal for any stabilizer code; for the seven-qubit code $H^{\otimes 7} = \bar{H}, P^{\dagger \otimes 7} = \bar{P}$, and transversal CNOT = $\bar{\text{CNOT}}$ (for any CSS code), giving the logical Clifford group transversally. But:

Theorem (Eastin–Knill). *No quantum code admits a universal set of transversal logical gates.*

9 Fault-tolerant error correction

Naive syndrome extraction lets one ancilla error spread to many data qubits. Three schemes fix this. **Shor EC** uses a verified n -qubit cat state $|0 \cdots 0\rangle + |1 \cdots 1\rangle$ so each data qubit touches only one ancilla qubit, and repeats the syndrome measurement for a consensus. **Steane EC** (CSS codes) prepares encoded $|\bar{0}\rangle + |\bar{1}\rangle$ and $|\bar{0}\rangle$ ancillas and uses transversal CNOT, so a single measurement fault stays a single bit. **Knill EC** teleports the encoded state into a fresh block, reading the syndrome off the encoded Bell measurement for any stabilizer code. Steane and Knill push effort into preparing *known*, offline-verifiable ancillas — the source of their high thresholds.

10 Completing universality: magic states

Since $R_{\pi/8}$ is not transversal it is supplied by **gate teleportation** through a special ancilla: teleporting then applying $\bar{U}$ equals teleporting through $(I \otimes \bar{U})(|\bar{00}\rangle + |\bar{11}\rangle)$ with correction $\bar{U}QU^\dagger$. This is favourable for $R_{\pi/8}$ because

$$R_{\pi/8} X R_{\pi/8}^\dagger = e^{i\pi/4} X P^\dagger \in \mathcal{P}_1 \cdot \mathcal{C}_1, \quad R_{\pi/8} Z R_{\pi/8}^\dagger = Z,$$

defining the **Clifford hierarchy** $\mathcal{C}_k = \{U : UPU^\dagger \subseteq \mathcal{C}_{k-1}\}$ with $R_{\pi/8} \in \mathcal{C}_3$. A transversal Clifford group plus a single verified magic state for $R_{\pi/8}$ gives a universal fault-tolerant set on the seven-qubit code; $\{H, \text{CNOT}, R_{\pi/8}\}$ is dense in $U(2^n)$ and **Solovay–Kitaev** approximates any gate to accuracy ϵ with $\text{polylog}(1/\epsilon)$ gates.

11 The threshold theorem

One studies **extended rectangles** (ExRecs): a gate gadget with its leading and trailing error corrections. An ExRec is *good* if it has at most t faults, and good implies correct; a circuit of only good ExRecs reproduces the ideal circuit. The payoff is **concatenation**: under a local stochastic error model the logical rate per level obeys

$$p' \leq A p^{t+1}, \quad p_T = 1/A^{1/t}, \quad \frac{p^{(j)}}{p_T} \leq \left(\frac{p}{p_T}\right)^{(t+1)^j},$$

so below threshold $p < p_T$ the logical rate falls doubly exponentially in the number of levels, with polylogarithmic overhead.

Theorem (Threshold theorem). *There is a threshold p_T such that for any $p < p_T$, any ideal circuit C can be simulated to within statistical distance ϵ by a fault-tolerant circuit using only $\text{polylog}(|C|/\epsilon)$ times as many qubits and time steps.*

Concrete thresholds depend on the code: a careful seven-qubit accounting gives $p_T \gtrsim 2.73 \times 10^{-5}$, the best rigorous bound is around 10^{-3} , and Knill’s teleportation-based schemes reach a few percent. The theorem assumes parallel operations, fresh ancilla preparation mid-computation, non-destructive errors, and bounded (local stochastic / non-Markovian) error correlations.

Takeaways

- No-cloning and continuous errors are evaded by linear encoding and digitization: correcting all weight- t Paulis corrects all t -qubit errors.

- The Knill–Laflamme conditions say when a subspace is a code; distance $d = 2t + 1$ and the Hamming / Singleton bounds constrain $[[n, k, d]]$.
- The stabilizer formalism turns codes into Abelian Pauli groups and syndromes into eigenvalue flips, linking quantum codes to classical and GF(4) coding theory.
- Fault tolerance tames error propagation; transversal gates are ideal but never universal.
- Universality is completed by teleporting magic states up the Clifford hierarchy.
- Below a constant threshold, concatenation suppresses logical error doubly exponentially with polylogarithmic overhead.

References

- [1] D. Gottesman, *An Introduction to Quantum Error Correction and Fault-Tolerant Quantum Computation*, arXiv:0904.2557 (2009).
- [2] P. W. Shor, *Scheme for reducing decoherence in quantum memory*, Phys. Rev. A **52**, 2493 (1995).
- [3] A. M. Steane, *Error correcting codes in quantum theory*, Phys. Rev. Lett. **77**, 793 (1996).
- [4] E. Knill and R. Laflamme, *A theory of quantum error-correcting codes*, Phys. Rev. A **55**, 900 (1997).
- [5] A. R. Calderbank and P. W. Shor, *Good quantum error-correcting codes exist*, Phys. Rev. A **54**, 1098 (1996).
- [6] B. Eastin and E. Knill, *Restrictions on transversal encoded quantum gate sets*, Phys. Rev. Lett. **102**, 110502 (2009).
- [7] P. Aliferis, D. Gottesman, and J. Preskill, *Quantum accuracy threshold for concatenated distance-3 codes*, Quant. Inf. Comp. **6**, 97 (2006).